

Basic Concepts of Artificial Intelligence

A beginner-friendly guide to how AI works,
the ideas behind it, and where it's used

Prepared as a quick-reference introduction to AI fundamentals

Contents

1. What Is Artificial Intelligence?
2. Machine Learning — Teaching Computers from Data
3. Types of Machine Learning
4. Neural Networks and Deep Learning
5. Key Terms You'll Keep Hearing
6. Categories of AI by Capability
7. Everyday Applications of AI
8. Limitations and Things to Keep in Mind

1. What Is Artificial Intelligence?

Artificial Intelligence (AI) is the branch of computer science focused on building systems that can perform tasks which normally require human intelligence — things like recognizing images, understanding language, making decisions, and solving problems.

Rather than being explicitly programmed with rules for every situation, many modern AI systems learn patterns from data and use those patterns to make predictions or decisions on new, unseen situations.

A simple way to think about it:

- **Traditional programming:** a human writes exact rules → the computer follows them.
- **AI / Machine Learning:** a human provides examples (data) → the computer finds the rules itself.

2. Machine Learning — Teaching Computers from Data

Machine Learning (ML) is the most widely used approach to AI today. It refers to algorithms that improve their performance on a task by learning from data, instead of following fixed instructions.

The basic ML workflow usually looks like this:

- **Collect data** — gather examples relevant to the task (e.g. photos, text, numbers).
- **Train a model** — feed the data into an algorithm so it can learn patterns.
- **Test/evaluate** — check how well the model performs on data it hasn't seen before.
- **Deploy and use** — apply the trained model to make predictions in the real world.

3. Types of Machine Learning

Type	How It Works	Example
Supervised Learning	Learns from labeled data (input paired with the correct answer)	Predicting exam scores from study hours
Unsupervised Learning	Finds hidden patterns or groups in unlabeled data	Grouping customers by shopping behavior
Reinforcement Learning	Learns by trial and error, receiving rewards or penalties	An AI learning to play chess by playing many games

4. Neural Networks and Deep Learning

A neural network is a type of ML model loosely inspired by how neurons in the brain connect and pass signals. It consists of layers of small computing units ("neurons") that each transform incoming data slightly and pass it to the next layer.

Deep Learning refers to neural networks with many layers ("deep" networks). More layers let a model learn increasingly abstract features — for example, in image recognition, early layers might detect edges, middle layers detect shapes, and later layers detect entire objects like faces.

Deep learning powers most of today's headline AI systems, including image recognition tools, voice assistants, and large language models like the ones behind modern chatbots.

5. Key Terms You'll Keep Hearing

Algorithm: A step-by-step procedure or set of rules a computer follows to solve a problem or perform a computation.

Model: The output of training — a system that has learned patterns from data and can make predictions.

Training Data: The dataset used to teach a model; the model learns patterns from these examples.

Dataset: An organized collection of data used for training or testing a model.

Parameters / Weights: Internal values a model adjusts during training to improve its predictions.

Overfitting: When a model learns the training data too closely, including its noise, and performs poorly on new data.

Natural Language Processing (NLP): The field of AI focused on understanding and generating human language.

Computer Vision: The field of AI focused on interpreting images and video.

Large Language Model (LLM): A deep learning model trained on huge amounts of text, able to understand and generate human-like language.

Bias (in AI): Systematic errors or unfair patterns a model can pick up from imbalanced or unrepresentative data.

6. Categories of AI by Capability

AI is sometimes classified by how broadly capable a system is:

- **Narrow AI (Weak AI):** designed for one specific task, such as spam filtering, facial recognition, or playing chess. This describes virtually all AI in use today.
- **General AI (Strong AI):** a hypothetical system with human-level intelligence across any task — this does not exist yet and remains a long-term research goal.
- **Superintelligence:** a theoretical future AI surpassing human intelligence in all domains; currently speculative and widely debated.

7. Everyday Applications of AI

- **Recommendation systems** — suggesting videos, songs, or products (e.g. streaming and shopping apps).
- **Virtual assistants** — voice assistants that understand spoken commands.
- **Search engines** — ranking and retrieving relevant web results.
- **Healthcare** — assisting in medical image analysis and disease prediction.
- **Navigation apps** — predicting traffic and suggesting optimal routes.
- **Translation tools** — converting text or speech between languages.
- **Chatbots** — answering questions and assisting with customer support.
- **Fraud detection** — spotting unusual patterns in banking transactions.

8. Limitations and Things to Keep in Mind

- **AI learns from data** — if the data is biased or incomplete, the model's outputs can be too.
- **Lack of true understanding** — most AI systems recognize patterns rather than "understand" meaning the way humans do.
- **Errors and hallucinations** — AI models can produce confident but incorrect outputs, so results usually need human review.
- **Data privacy** — training and using AI often involves large amounts of personal or sensitive data.
- **Explainability** — complex models (especially deep learning) can be hard to interpret, making it difficult to know exactly why a decision was made.